

Chapter 2: Case study

Introduction

In this chapter we will look at a simple graph problem and see how to solve it so that we are using close to 100% of the theoretical maximum performance of the computer.

The shortcut problem

We will use the following simple graph problem as our running example throughout this chapter. We have a directed graph with n nodes. The nodes of the graph are labeled with numbers $0, 1, \dots, n - 1$. There is a directed edge between each pair of the nodes; the **cost** of the edge between nodes i and j is d_{ij} ; we will assume that d_{ij} is a nonnegative real number. For convenience, we write $d_{ii} = 0$ for each node i .

However, the costs do not necessarily satisfy the triangle inequality. We might have an edge of cost $d_{ij} = 10$ from node i to j , but there might be an intermediate node k with $d_{ik} = 2$ and $d_{kj} = 3$. Then we can follow the route $i \rightarrow k \rightarrow j$ at a total cost of $2 + 3 = 5$, while the direct route $i \rightarrow j$ would cost 10.

Our task is to find for all i and j what is the cost of getting from i to j by taking **at most two edges**. If we write r_{ij} for the result, then clearly

$$r_{ij} = \min_k (d_{ik} + d_{kj}),$$

where k ranges over $0, 1, \dots, n - 1$. Note that we will also consider here e.g. the route $i \rightarrow i \rightarrow j$, and hence we will also find a path of one edge if it happens to be cheapest.

Remarks [advanced]

We can write the distances d_{ij} as a matrix D , and the results r_{ij} as a matrix R . Then $R = D \circ D$, where \circ denotes **min-plus matrix multiplication**.

If we iterate the process and calculate

$$D_2 = D \circ D, \quad D_4 = D_2 \circ D_2, \quad D_8 = D_4 \circ D_4, \quad \dots$$

we can quickly find the cost of getting from any node to any other node by following at most 2, 4, 8, ... edges. Iterating this operation for a logarithmic number of times would therefore also give **all-pairs shortest path distances** (but recall that there are **more direct ways** of solving the all-pairs shortest path problem).

Interface

We will implement a C++ function `step` with the following prototype:

```
void step(float* r, const float* d, int n);
```

Here `n` denotes the number of nodes, `d` contains the input, and `r` will contain the result. Both `d` and `r` are pointers to arrays with `n * n` floats.

The cost of the edge from node `i` to node `j` is stored in `d[n*i + j]`. Similarly, the cost of getting from `i` to `j` by following at most two edges will be stored in `r[n*i + j]`. Here `0 <= i < n` and `0 <= j < n`.

Example

Here is a simple example of how we could call `step`:

```
int main() {
    constexpr int n = 3;
    const float d[n*n] = {
        0, 8, 2,
        1, 0, 9,
        4, 5, 0,
    };
    float r[n*n];
    step(r, d, n);
    for (int i = 0; i < n; ++i) {
        for (int j = 0; j < n; ++j) {
            std::cout << r[i*n + j] << " ";
        }
        std::cout << "\n";
    }
}
```

This program should produce the following output:

```
0 7 2
1 0 3
4 5 0
```

Here, for example, there is a direct edge from node 0 to node 1 with cost 8. However, we can take an edge from node 0 to node 2 (cost 2), and then an edge from node 2 to node 1 (cost 5), which results in a path of total cost 7.

