


Chapter 3: Multithreading with OpenMP

Parallelizing nested loops

If we have nested for loops, it is often enough to simply **parallelize the outermost loop**:

```
a();
#pragma omp parallel for
for (int i = 0; i < 4; ++i) {
    for (int j = 0; j < 4; ++j) {
        c(i, j);
    }
}
z();
```




This is all that we need most of the time. You can safely stop reading this part now; in what follows we will just discuss what to do in some rare corner cases.

Challenges


Sometimes the outermost loop is so short that not all threads are utilized:

```
a();
#pragma omp parallel for
for (int i = 0; i < 3; ++i) {
    for (int j = 0; j < 6; ++j) {
        c(i, j);
    }
}
z();
```



We could try to parallelize the inner loop. However, then we will have more overhead in the inner loop, which is more performance-critical, and there is no guarantee that the thread utilization is any better:

```
a();
for (int i = 0; i < 3; ++i) {
    #pragma omp parallel for
    for (int j = 0; j < 6; ++j) {
        c(i, j);
    }
}
z();
```



Good ways to do it


In essence, we have got here $3 \times 6 = 18$ units of work, and we would like to spread it evenly among the threads. The correct solution is to **collapse it into one loop** that does 18 iterations. We can do it manually:

```
a();
#pragma omp parallel for
for (int ij = 0; ij < 3 * 6; ++ij) {
    c(ij / 6, ij % 6);
}
z();
```



Or we can ask OpenMP to do it for us:

```
a();
#pragma omp parallel for collapse(2)
for (int i = 0; i < 3; ++i) {
    for (int j = 0; j < 6; ++j) {
        c(i, j);
    }
}
z();
```



Either of the above solutions are just fine.

Wrong way to do it, part 1

Unfortunately, one often sees failed attempts of parallelizing nested for loops. This is perhaps the most common version:

```
a();
#pragma omp parallel for
for (int i = 0; i < 3; ++i) {
    #pragma omp parallel for
    for (int j = 0; j < 6; ++j) {
        c(i, j);
    }
}
z();
```

This code **does not do anything meaningful**. "Nested parallelism" is disabled in OpenMP by default, and the second pragma is ignored at runtime: a thread enters the inner parallel region, a team of only one thread is created, and each inner loop is processed by a team of one thread. The end result will look, in essence, identical to what we would get without the second pragma — but there is just more overhead in the inner loop:



On the other hand, if we tried to enable "nested parallelism", things would get much worse. The inner parallel region would create more threads, and overall we would have $3 \times 4 = 12$ threads competing for the resources of 4 CPU cores — not what we want in a performance-critical application.

Wrong way to do it, part 2

One also occasionally sees attempts of using multiple nested `omp for` directives inside one `parallel` region. This is **seriously broken**; OpenMP specification does not define what this would mean but simply forbids it:

```
a();
#pragma omp parallel for
for (int i = 0; i < 3; ++i) {
    #pragma omp for
    for (int j = 0; j < 6; ++j) {
        c(i, j);
    }
}
z();
```

In the system that we have been using here as an example, the above code thankfully gives a **compilation error**. However, if we manage to trick the compiler to compile this, e.g. by hiding the second `omp for` directives inside another function, it turns out that the program **freezes** when we try to run it.